

Software segmentation offloading for FreeBSD

Stefano Garzarella, Luigi Rizzo
stefano.garzarella@gmail.com



Università di Pisa

September 28, 2014



Outline

- Introduction
- Hardware offloading (TSO, RSC)
- Software offloading (LRO)
- Generic Segmentation Offload (GSO)
- Results




Introduction

- The use of large frames makes network communication much less demanding for the CPU.
- Backward compatibility and slow links requires the use of **1500** byte or smaller frames.
 - Modern NICs with hardware TCP segmentation offloading (**TSO**) address this problem.
 - Generic software version (**GSO**) provided by the OS has reason to exist, for use on paths with no suitable hardware.


Hardware Offloading

Sender side

- TCP Segmentation Offload (**TSO**)
 - Data segmentation is offloaded to the NIC, that divides the data into the maximum transmission unit (MTU) size of the outgoing interface.
 - The network stack is crossed only once per (large) segment instead of once per 1500-byte frames.
 - Issues:
 - Only TCP traffic
 - Early version supported only IPv4




TCP Segmentation Offload




Hardware Offloading

Receiver side

- Receive Side Coalescing (**RSC**)
or hardware LRO
- Allows a NIC to combine incoming TCP/IP packets that belong to the same connection into one large receive segment before passing it to the operating system.
- Reduces CPU use because the TCP/IP stack is executed only once for a set of received Ethernet packets.



Receive Side Coalescing






Software offloading

- Why software implementation?
 - Older NICs (IPv4 / IPv6)
 - Buggy NICs
 - Communication between Virtual Machines
 - Easy to extend for new protocols



Software offloading

Receiver side



- Large Receive Offload (**LRO**)
 - Software implementation of RSC
 - Available since FreeBSD 7.1
 - Requires changes in each device drivers


Large Receive Offload




Software offloading

Sender side

- **Generic Segmentation Offload (GSO)**
 - Software implementation of TSO
 - Supported TCP, UDP and IPv4, IPv6
 - Available for FreeBSD:
 - -current, 10-stable, 9-stable
 - Segmentation just before the device driver




GSO on TCP flow



GSO: data flow

- Network stack changes
 - `tcp_output()`
 - `ip_output()`
 - `ether_output()`
 - `gso_dispatch()`



tcp_output()


1. Checks if GSO is enabled:

- `sysctl net.inet.tcp.gso = 1`
- `sysctl net.gso."ifname".enable_gso = 1`

2. Checks if the packet length exceeds the MTU


- If 1 and 2 are true, sets GSO flag

- `m->m_pkthdr.csum_flags |= GSO_TO_CSUM(GSO_TCP4);`




ip_output()

- If GSO is enabled and required:
 - avoids checksum (IP & TCP)
 - avoids IP Fragmentation



ether_output()


- If GSO is enabled and required:
 - calls `gso_dispatch()` instead of `ifp->transmit()`





gso_dispatch()

```
enum gso_type {      int (*gso_functions[GSO_END_OF_TYPE])  
                     (struct ifnet*, struct mbuf*, u_int);  
    GSO_NONE,  
    GSO_TCP4,  
    GSO_TCP6,  
    #define CSUM_TO_GSO(x) (((x & CSUM_GSO_MASK)  
    GSO_UDP4,           >> CSUM_GSO_OFFSET)  
    GSO_UDP6,  
/*  
 * GSO_SCTP4, TODO  
 * GSO_SCTP6,  
 */  
    GSO_END_OF_TYPE  
};  
  
int  
gso_dispatch(struct ifnet *ifp,  
             struct mbuf *m, u_int mac_hlen)  
{  
    ...  
    gso_flags = CSUM_TO_GSO(m->m_pkthdr.csum_flags);  
    ...  
  
    error = gso_functions[gso_flags](ifp, m, mac_hlen);  
    return error;  
}
```






gso_functions[]

- gso_functions[GSO_TCP4]
 - gso_ip4_tcp(...) - GSO on TCP/IPv4 packet
 1. `m_seg(struct mbuf *m0, int hdr_len, int mss, ...)`
returns the mbuf queue that contains the segments of the original packet (m0).
 - `hdr_len` - first bytes of m0 that are copied in each new segments
 - `mss` - maximum segment size
 2. fixes TCP and IP headers in each new segments
 3. sends new segments to the device driver [`ifp->if_transmit()`]




GSO: m_seg()







GSO: fix TCP/IPv4 headers



GSO: fix TCP/IPv6 headers




GSO on UDP flow






GSO: fix UDP/IPv4 headers



GSO: fix UDP/IPv6 headers





sysctl

To manage the GSO parameters there are some **sysctl**:

- **net.inet.tcp.gso**
GSO enable on **TCP** communications (!=0)
- **net.inet.udp.gso**
GSO enable on **UDP** communications (!=0)
- for each interface:
 - **net.gso.dev."ifname".max_burst**
GSO burst length limit [default: IP_MAXPACKET=65535]
 - **net.gso.dev."ifname".enable_gso**
GSO enable on "ifname" interface (!=0)



GSO code

- Kernel patches for FreeBSD-current, FreeBSD 10-stable and FreeBSD 9-stable available at:
<https://github.com/stefano-garzarella/freebsd-gso>
- FreeBSD source with GSO available at:
<https://github.com/stefano-garzarella/freebsd-gso-src>
- To compile kernel with GSO support:
 - “**options GSO**” in kernel config



GSO patch

diff gso-current


sys/conf/NOTES		2 +
sys/conf/files		1 +
sys/conf/options		1 +
sys/net/gso.c		1020 ++++++
sys/net/gso.h		174 +++++++
sys/net/if_ETHERSUBR.C		31 ++
sys/net/if_var.h		1 +
sys/netinet/ip_output.c		43 +-
sys/netinet/tcp_input.c		11 +
sys/netinet/tcp_output.c		84 +++-
sys/netinet/tcp_subr.c		15 +
sys/netinet/tcp_var.h		6 +-
sys/netinet/udp_usrreq.c		22 +-
sys/netinet/udp_var.h		6 +
sys/netinet6/ip6_output.c		51 +-
sys/netinet6/udp6_usrreq.c		8 +
sys/sys/mbuf.h		2 +
17 files changed, 1459 insertions(+), 19 deletions(-)		



Experiments


- Sender: CPU i7-870 at 2.93 GHz + Turboboost, Intel 10 Gbit NIC.
- Receiver: CPU i7-3770K at 3.50GHz + Turboboost, Intel 10 Gbit NIC.
 - (RSC/LRO)-enabled (otherwise TSO/GSO are ineffective)
- Benchmark tool: netperf 2.6.0

Results TCP/IPv4




Freq. [GHz]	TSO	GSO	none	Speedup GSO - none
2.93	9347	9298	8308	12%
2.53	9266	9401	6771	39%
2.00	9408	9294	5499	69%
1.46	9408	8087	4075	98%
1.05	9408	5673	2884	97%
0.45	6760	2206	1244	77%

Results TCP/IPv6




Freq. [GHz]	TSO	GSO	none	Speedup GSO - none
2.93	9097	8861	4966	78%
2.53	9113	8290	4008	107%
2.00	9066	6599	3152	109%
1.46	7357	5180	2348	121%
1.05	6125	3607	1732	108%
0.45	2005	1505	651	131%

Results UDP/IPv4



Freq. [GHz]	GSO	none	Speedup GSO - none
2.93	9440	8084	17%
2.53	7772	6649	17%
2.00	6336	5338	19%
1.46	4748	4014	18%
1.05	3359	2831	19%
0.45	1312	1120	17%

Results UDP/IPv6



Freq. [GHz]	GSO	none	Speedup GSO - none
2.93	7281	6197	17%
2.53	5953	5020	19%
2.00	4804	4048	19%
1.46	3582	3004	19%
1.05	2512	2092	20%
0.45	998	826	21%



Future works

- More performance measurements
- Optimize code paths
- Add support to new protocols (SCTP, ...)



Thank you!